


Developing a StarCraft: Brood War Agent

Carsten Nielsen

DTU



Kongens Lyngby 2015

Technical University of Denmark
Department of Applied Mathematics and Computer Science
Richard Petersens Plads, building 324,
2800 Kongens Lyngby, Denmark
Phone +45 4525 3031
compute@compute.dtu.dk
www.compute.dtu.dk

Summary (English)

The goal of the thesis is to ...

Summary (Danish)

Målet for denne afhandling er at ...

Preface

This thesis was prepared at DTU Compute in fulfillment of the requirements for acquiring an M.Sc. in Engineering.

Starcraft is one of the strongest game franchise ever created. Professional players compete in Starcraft 2 tournaments with large price purses. Perhaps less known is that former developers of Othello and chess playing programs have taken on the task to develop AIs for the game. Most developers use the relatively old Starcraft Brood War as all the necessary infrastructure is readily available: <http://code.google.com/p/bwapi/>. Championships for AIs, bots, are arranged yearly, see for example <http://www.sscaitournament.com/> and <http://webdocs.cs.ualberta.ca/~cdavid/starcrafttaicomp/>. In addition, ladder systems are also available, see <http://bots-stats.krasi0.com/>.

The project focuses on developing a new bot and is aimed at 20 ECTS.

The project involves:

- studying and implementing AI methods for handling large search spaces.
- developing high-level strategy reasoning algorithms and heuristics.
- solving unstructured optimization problems.

Lyngby, 30-June-2015

Not Real

Carsten Nielsen

Acknowledgements

I would like to thank my...

Contents

Summary (English)	i
Summary (Danish)	iii
Preface	v
Acknowledgements	vii
1 Introduction	1
2 Project Process	3
2.1 Testing	3
3 StarCraft	5
4 Related Works	7
4.1 Strategies in StarCraft	7
4.2 Bots in StarCraft	8
5 Agent Design	11
5.1 Overall Strategy	11
5.2 Overall Architecture Design	12
6 Workers and Resources	15
6.1 Managing Workers	16
6.2 Mining Minerals	17
6.3 Harvesting Gas	18

7	Production	19
7.1	Production Architecture	20
7.2	Training Units	20
7.3	Building Structures	22
7.4	Building Expansions	25
8	Exploration	27
8.1	Tracking Units	28
8.2	Scouting	29
9	Combat	31
9.1	Attacking	31
9.2	Defending	35
10	Macro Strategy	37
10.1	Economy	38
10.2	Combat	39
10.3	Build-orders	40
11	Results	43
12	Conclusion	45
	Bibliography	47

CHAPTER 1

Introduction

Artificial Intelligence in *real-time strategy* (RTS) games offer the challenge of limitless decision space, incomplete world information and varied strategies in a shifting *metagame*. A computer controlled player, called a *bot*, must both be efficient and effective. Compared to other games, RTS focuses on a single agent.

StarCraft: Brood War, released by Blizzard in 1998, is one such RTS game that has been the focus of research. In addition to usual RTS elements, StarCraft is asymmetric with three different factions to play as and slightly asymmetric maps. Since the advent of the unofficial *Brood War API* (BWAPI), people have been able to develop their own bots for the game, pitting them against each other and spawning a few tournaments. The tournament scene is still maintained even though both the game and API have become dated.

This project will both focus on creating a competitive bot by following successful contemporary bots, and investigate how prevailing bots are designed and how they overcome the aforementioned challenges. The bot will be spar against other bots in the SSCAI (Student StarCraft AI) tournament, possibly competing in the 2015 tournament (depending on when it starts).

Original Project Plan

Original project plan.

Revised Project Plan

Revised project plan.

Brief self-evaluation.

CHAPTER 2

Project Process

The bot was developed with the agile development process. The implementation and design was iterated upon during short week-long sprints. After each there would be an evaluation of the bot's current status and new goals would be identified for the next sprint. A backlog kept during development was consulted when determining the status and goals in scope of the whole time-plan. There were a few set milestones throughout development, usually a month apart. They determined some requirements the bot should satisfy in relation to its performance at specific points in the development. The milestones served as a medium between sprints and the overall goal, ensuring the bot improved and did so consistently.

The first milestone were attaining the minimum viable product. The requirements for this was a bot capable of defeating a passive opponent, implying means to gather resources, build units and attacking. This would imply the bot had a non-zero chance at winning a round, making it an actual competitor.

2.1 Testing

Testing was done in two steps. First the robustness of the implementations was tested by playing the bot with random built-in StarCraft opponents and

tournament maps. This was to ensure the stability of the bot and that the implementations worked as required. The second step was uploading the bot to the SCCAIT tournament page, where the bot would play against the other official bots. This would serve to test the strategic integrity of the bot, and was done rarer than the first step. Because it is more difficult to retrieve data from rounds on the SCCAIT, the first step was needed to determine crash causes and such technicalities.

The opponents on the SSCAIT should be at least updated since that last tournament, and would therefore be viable candidates to test the bot's strength. It is required to make ones bot open source if it is to compete, so there is no reason to keep the newest bot version a secret beyond a year.

CHAPTER 3

StarCraft

StarCraft is a *real-time strategy* game released by Blizzard Entertainment in 1998. Players take the role as one of three different factions, the *Terran*, *Zerg* or *Protoss*. Controlling structures and troops, they seek to expand their economy, technology and army until they can eliminate all opponents. The expansion *StarCraft: Brood War*, also released in 1998, increased the roster of units available among other things.

In the multiplayer mode, players start with a single resource depot structure and a few worker units. They must from there destroy all the opponents' structures. The matches happen on an array of different *maps*. Players must deal with imperfect map information as their visibility is limited to their units and their immediate surroundings - the rest is shrouded in the *fog-of-war*. With a large amount of different units and buildings available, the players compete by adapting strategies to counter their opponents. There is no perfect strategy capable of winning every game, as the decision space is high infinite.

CHAPTER 4

Related Works

Bots in StarCraft

4.1 Strategies in StarCraft

In most RTS games, StarCraft being no exception, there is a trinity of strategies in a rock-paper-scissor formation. These are the *rush*, *boom* and *turtle* strategies.

The *rush* strategy is when the player attempts to produce and attack with troops as fast as possible. There are many variations in StarCraft, where the most aggressive are called *all-in*, as the rushing player will lose shortly afterwards if the initial rush fails. This could happen if the opponent builds stationary defenses, which are usually much more cost-effective than the mobile units required to attack.

Closely related, a player could instead focus on developing economy or technology to gain advantage against slow strategies. This is usually called a *boom* strategy, which in StarCraft involves gaining an *economic advantage*. If successful, the player will attain superior units and easily replace lost ones, winning either through strength or attrition. The weakness is early game defense will have to be sacrificed which puts the boom player at risk to enemy rush strategies.

Finally, the defensive strategy called a *turtle* or *turtling* builds strong defenses early on. As mentioned, this will counter an enemy rush by costing fewer resources, used to expand in economy or technology. The resulting advantage is then used to win. This is a slow strategy however, which fails against the riskier *boom* strategy. Since the defensive strategy does not put pressure on the opponent, they are free to completely spend resources on expanding to stronger units.

A very important part of StarCraft strategies are the openings. Much like chess openings, StarCraft has a wide array of detailed steps to take in the first few minutes of a round.

Usually either of these strategies are used during the early game. Later on, players will have more resources to pursue multiple goals and a more balanced overall strategy. This takes form in StarCraft by players attempting to gain a long-term economic advantage, done by controlling the limited resource locations. While early game can be decided by a single battle, later stages require multiple.

4.2 Bots in StarCraft

Usually bots focus on only one strategy. Some of the more advanced bots identify whatever opening strategy the opponent is using.

Few bots reach the late-game content of StarCraft. Those that do usually focus on the incredible precision and speed which the bot can play compared to a human.

API

BWAPI is an API for StarCraft Broodwar and is injected upon startup by *ChaosLauncher*. It loads an AI module written in c++, which can retrieve information about the current match's map status and send commands to the game through BWAPI. This controls the player's actions and allows the development of agents for StarCraft.

Depending on the settings in BWAPI, the agent can be disallowed to retrieve information a player would not be able to. This means some units are in some levels of accessibility, where invisible and destroyed units are completely inac-

cessible. This has some limits however, as the agent can retrieve information about burrowed and cloaked units as if they were not. While a keen player can spot these units as they are not completely transparent, the agent has no limits as if they were completely visible, giving it an advantage. Additionally the agent is not limited to what is visible currently on the display, but can retrieve information from all over the map.

There is a popular library, the *Broodwar Terrain Analyzer* or *BWTA* for short, which pre-processes maps for locations of interest. Most importantly, it marks the optimal depot locations for harvesting resources. While BWAPI already does this for start-locations, BWTA also does this for all resource clusters, marking viable expansion locations. The pre-processing of the map only has to be done once as the results are stored for subsequent games on the map. This library is automatically included in the 3.7.4 BWAPI test bot, so it is probable that almost all bots use the library.

As an additional note, BWAPI divides the map into *regions*. These contain satellite data about the borders to neighboring regions, included resources and base locations. They are important because the regions are separated in *choke-points*, which are either ramps or bottlenecked passages. Usually a region only contains a single base location and resource cluster, however some regions contain multiple clusters with overlapping base locations. Because of these qualities regions can be used to mark occupied locations by players.

While newer versions of BWAPI exist, the 3.7.4 version is still widely used as it is considered the most stable, and is compatible with BWTA. There does exist a BWTA clone for 4th versions, however it is not very stable. The downside to 3.7.4 is that it does not support c++ 11. There does not exist an equivalent to BWAPI for StarCraft 2, since technically BWAPI is a hack and using it would result in a ban. AI's have been scripted in StarCraft 2's own editor however, and some have even gone and interpreted the display output of the game for a bot.

AI Competitions

Different competitions and their descriptions.

Competition rules: Setup, limits, goals.

Ties impossible

This bots involvement/predicted involvement.

CHAPTER 5

Agent Design

In this chapter the agents macro-strategy will be described as well as its overall architecture and implementation design.

TODO Intro.

5.1 Overall Strategy

Before any strategies for mid-game or late-game could be devised, the bot must first master early-game. If the bot loses in early-game, the match will never proceed to later stages. Compared to matches between humans, the game usually ends much quicker, rarely leaving early-game. This is because later stages of the game involves many more kinds of units and strategies, such that the AI must be more advanced. Therefore, it seems prudent to first perfect the early game of the bot. From there, if time allows, the bot could attempt mid-game tactics.

Simplest solution is implementing one strategy, which can be expanded into mid-game if the round drags on that long. Both boom and turtle strategies seek to push the match out of early game, so the remaining strategy which lies in

early game is rush. This has seen a lot of use in the bot tournaments, probably being the most popular choice. It requires nothing but building the earliest combat unit and attacking. When moving on to mid-game, additions could include upgrading, expanding to new resources or moving on to more advanced units. Even a failed rush would put pressure on the opponent, allowing the bot to compete in later stages of the match. This makes such a strategy sustainable in case the bot reaches mid-game, but is also proven to be a strong early-game strategy.

The bot will only focus on a single race. Even though Protoss are the slowest of the three, they also have the strongest and easiest to use early troops. This makes it a strong rush candidate, as it can easily beat opposing races' rushes if prepared. While the main AI challenges between races are shared, gameplay elements differ quite a lot.

5.2 Overall Architecture Design

The imperative when designing the bot is to lessen the decision space as greatly as possible. The problems the bot faces are easily divided into smaller, isolated problems. Its therefore possible to construct the bot out of individual *modules*, each solving some subset of problems. This isolation lessens the amount of concurrent decisions and information available to the agent, reducing the decision space.

These modules are structured in a *hierarchy*. A superior module can command a subordinate, however only in terms of a *black box*, without knowledge of its internal structure. By limiting the information available to the superior, we can easily limit the decision space with clever design. The clear benefit of this structure is easy replacement of individual modules, without damaging or refactoring neighbors. On the other hand, modules will inherently be limited with information, but we must trust that not all available data is required or even relevant for all decisions.

The *UAlberta* bot by Dave Churchill uses such a module structure, however it places them in an interesting hierarchy. The modules in the bot are in a *arborescence* graph structure, that is, there is a root module which has exactly one path to each other node in the hierarchy. In other words there are no cyclic dependencies, and it is related to a tree structure, except a single module can have multiple parents. Churchill alleges that it is based on "proven military structures" - in any case the UAlberta bot is high-ranking bot, victor of one competition and runner up in others. The datastructure must have been proven

to work by now.

The benefits of this structure is lesser and easier dependencies, removing the inherent challenges in cyclic dependencies. It enhances the benefits of modular design, since modules has a stricter position in the hierarchy, making the modular design easier. It is also a great boon to agile development, as the tree will simply evolve upwards with newer modules as superiors to the older. The lower modules make smaller decisions with fewer resource such as building and gathering, and the higher modules control more information to make the larger scale decisions such as strategies and attacking.

On the other hand, the structure allows less complicated interactions. By limiting the hierarchy, the information available is limited as some modules must be at the bottom of the hierarchy. These are inevitably void of interactions with higher modules. This is also felt at the root of the structure, as at some point all the choices have to be made in the final module.

TODO Hierarchy diagram + description.

In the following chapters, we will describe the different groups of modules in the hierarchy and the problems they solve. While the modules are individually separate, some clusters of them are also separate. These can be separated in chapters as few of the modules has any connections between. The order of clusters is bottom-up so as to follow the chain of command. These groups are *resources*, *production*, *information*, *combat* and *strategy*.

CHAPTER 6

Workers and Resources

This chapter covers the basics of gathering resources in StarCraft with a bot.

In early- and mid-game, *economic advantage* is the most important aspect of StarCraft strategy. This occurs when you have a higher resource intake than your opponent, such as if you have more workers that are gathering or more bases to gather from. The advantage is while maintaining an equivalent army to the opponent, one can replace lost troops faster, gain a larger army or upgrade the current one. If you eclipse your opponent in resources, you can perform worse in combat and still win. Without the advantage, upgrading troops would slow troop production down and the opponent could produce an equivalent or even stronger army at any time.

There are three kinds resources in StarCraft. *Workers* are the only units capable of gathering resources. Each race has one worker type and each player start the game with a couple of their race's type. Players also start with a *resource depot* or simply *depot*. This structure can produce new workers and is also the drop off point for gathering resources. The Protoss worker is the *probe* and their depot is the *nexus*.

Minerals are mined from mineral fields which usually are in clusters of 8-12. All units cost minerals, so it is the most important resource. Workers that mine these bring them back to the depot in batches of six mineral units. Depot's are

limited in their proximity to the clusters, but at optimal distance there can be 3 workers on each field. Mineral fields are usually placed in a half-circle formation, such that there exists an optimal depot location. While the first worker on each mineral field will gather at a linear rate, they will yield diminishing returns on the second and especially third worker.

Vespene Gas or simply *gas* is harvested from *refineries*. Each race has a distinct refinery structure, although they are very similar. These must be built upon *Vespene Geysers* of which there are up to two by each mineral cluster. It is harvested in batches of eight with a max of three workers per refinery at optimal depot distance. Advanced structures, units and all technologies depend on gas. The immediate cost of the refinery and lost mineral gathering is a liability against fast openings, but harvesting too late means fighting against stronger units.

The last resource, *supply* is a population limit and is not gathered like the other resources. Each unit reserves some amount of supply which is released upon their destruction. The only way to secure more supply is building the race's *supply structure*. Each of these add some amount of supply to the total, which is subtracted if they are destroyed. Units cannot be produced if they require more supply than the player has free. If the player happens to reserve more supply than their total, they will be unable to build more units until more supply is acquired, called *supply blocked*. Although supply is specifically the Terran resource, it is used as the general term for all races. *Psi* and *control* are the Protoss and Zerg equivalents. The Protoss supply structure is the *pylon*.

While multiple workers can gather from the same field or refinery, only one worker can actively occupy it. The rest is either returning cargo, moving to the resource or waiting.

In the next sections we will first describe the architecture behind worker management, followed by how minerals is mined and gas is harvested. Building supply is first covered in chapter 10, as this resources behaves much different than the others.

6.1 Managing Workers

The *task master* module contains workers related to a single depot. Every depot has its own instance of the module, separating the workers into resource clusters for easy gathering. The task master keeps the workers in sets and also designates a *task* for each worker. Initially a worker is tasked as *idle*, but other modules

could re-task workers, essentially allocating and freeing workers. Each task has a related set of workers, stored as a dictionary of sets. Given n workers and a small amount of tasks, insertions, deletions and re-tasks operations has time complexity $O(\log n)$. We expect n to be in the range of 0 – 24, so the operations are quite fast. The current set of tasks are **idle**, **mine**, **harvest**, **build** and **defend**.

Every task master is wrapped in a *vassal* module, all of which are contained in a dictionary within the *landlord* module. As the amount of needed vassals are unknown, they are stored in dynamic memory. The landlord designates new workers to or destroyed from the related Vassal. This allows for easier operations from superior modules to the worker pool.

None of these three modules have any agency, only responding to operations from superior modules or events from BWAPI.

The *gatherer* module commands all gathering workers, superior only to the landlord module and its immediate subordinates. Resource gathering was initially kept within each task master, however this makes it difficult to control gathering on a large scale. The resource priorities should not necessarily be uniform across all bases. Every frame, the Gatherer iterates through the vassals, commanding current gatherers and re-tasking all idle workers to gathering. How this is handled is explained the following sections.

6.2 Mining Minerals

StarCraft has a built-in worker gathering AI to help human players. Workers will automatically return cargo from resources (unless they are interrupted) and return to the same resource afterwards. When gathering minerals, they will move to another mineral field if the current one is occupied. This is an inefficient solution, as the worker could be stuck moving between minerals for long periods of time. It is also a liability for a bot, as the AI cannot be disabled. When harvesting gas they will wait until the refinery is unoccupied.

The simplest mineral gathering implementation is ordering idle gatherers to mine some arbitrary mineral. At some point, the built-in AI will ensure the workers are optimally scattered. It will however not be scattered immediately, and some workers will be very inefficient while moving from mineral to mineral. A simple but effective addition would be to scatter the initial workers.

By maintaining a queue of minerals, we can optimally scatter the workers. The

first element of the queue is the mineral with fewest workers and the one in the back has the most. By continually assigning new workers to the first element and moving it to the back, we maintain a queue where the last mineral has at most one more worker than the first. Removing workers however requires finding the mineral in the queue, and should be done sparingly. Therefore, it is assumed any building or defending activity will be short and temporary, and workers assigned such will not be removed from the scattering. This might result in an ineffective scattering at some points. It is not clear however if optimizing the scattering at all times results in optimal resource output, as workers might be moved between minerals too often, resulting in less time mining.

If we maintain a dictionary of workers and their targets, we can assign new ones in constant time and retrieving targets in logarithmic time. This could be improved to amortized constant time with hashing. Removing a worker however requires a search through the queue which is linear time.

6.3 Harvesting Gas

Gas harvest is very much like mineral mining, but simpler. Rarely will players have fewer than three workers on each refinery, as refineries will be built only when needed. Furthermore, almost all late-game units require gas, and harvesting is slower than mining.

The implementation is identical to mineral mining. No AI tournament maps contain more than one refinery, but the gas harvesting still uses a priority queue. In this case however, all operations become constant time, so performance is not harmed while code is reused.

CHAPTER 7

Production

This chapter concerns the creation of units including the architecture for its management. First we explain the StarCraft terms of units and their production.

A *Unit* is any player controlled entity, structure or not. Resources are also units, controlled by a neutral "player", but are both immobile and indestructible. Some spell effects, like the Terran nuke, are oddly also classified as units but they are neither selectable nor controllable. The term unit rarely includes these however.

Production here refers to the player controlled act of creating a new unit. Within StarCraft, producing a structure is called *building* and *training* if it is a non-structure. Once a unit is created, there is a duration where it is *constructing* before becoming *complete*. A unit cannot execute player commands and has no abilities during construction. Constructing structure are placed in the world when they are built, while non-structure are hidden within their constructor. When a non-structure is complete, it appears at the nearest free space around its constructor. A unit consumes resources upon production, including supply requirements. A constructing non-structure can be canceled for a full refund.

BWAPI notifies the AI module whenever a unit is created or completed in separate events, which can be used to update internal data-structures.

In the following sections we first briefly describe our location in the hierarchy. Then we begin with training of non-structure, moving on to building structures. Finally we cover expansions and expanding.

7.1 Production Architecture

The *Accountant* module is used for keeping internal track of spent resources and scheduled units. Commanding a unit to build or train will not spend resources until at least the next frame. Therefore it is required to keep internal records of these, otherwise different modules might spend the same resources twice. Unit scheduling is useful in later frames, such that modules do not request the same unit multiple times.

Training is handled by the *Recruiter* module, while building is done by the *Architect*. These are separate modules as the two jobs are very different in implementation. While the Recruiter is quite low in the hierarchy, the Architect is not since it requires workers from the Landlord module. Both are superior to the Accountant.

There is a third module, the *Morpher*, which handles the rare case of a unit that executes a *morph* command. A morphing unit transforms from one type to another, which is most often seen in case of building refineries. Counter-intuitively, the geyser morphs into a refinery and changes ownership to the constructing player, but the refinery is still built by a worker with a build command. The Morpher exclusively monitors morphing units, while the Architect handles the building.

7.2 Training Units

The *Recruiter* needs to be low in the hierarchy, as it needs to be accessed by many other modules. It does not contain an independent AI, acting only through certain events and method calls. The current implementation is generic such that it can handle any train-able unit in the game.

To train units, the steps are twofold: issue the train command to a relevant, available trainer and then monitor the construction.

Commanding Trainer

Some units can train others, such as the depots that train workers. It is almost exclusively structures that can train. A unit can only train one at a time, although multiple can be queued up. Every unit is trained by at most one other unit, so a graph of trainers and trainees would be a *forest* structure.

Notice that it is inefficient to queue training, as this will lock resources. Instead, issuing the train command whenever the trainer is available will keep resources free without resorting to cancellation. It would be better to implement an internal queue if one was needed.

To command the training we require the trainer. For this we would like to keep records of all units capable of training. BWAPI inherently contains the trainer of any given unit type, so using a dictionary of trainers with their types as keys is sufficient and allows fast queries. From this we can search through the set of trainers until we find one that is available and then send the train command. The evaluation of trainers involve verifying the current existence and control of the unit, and ensuring it is not currently training another unit or has been commanded to in this frame.

Searching through the dictionary is logarithmic and iteration through the matching trainers is linear. Evaluating a trainer is constant. Given n trainers and m trainer matches, the time complexity for training is $\log(n) + m$, which could be improved to m with hashing. Neither n or m are usually very large, but m especially is only in the range of zero to four. Insertions and deletions of trainers are both logarithmic to their size.

Monitoring Trainee

The agent needs to know which units are currently scheduled, such that it can include this into its plans. Units could also be destroyed before completion, in which case the agent might need to reschedule.

Once the train command has been issued, the Accountant is notified of the costs and the type.

Upon receiving the event that a unit has been created, the Recruiter frees the resource costs from the Accountant, as the game state by now has withdrawn the costs itself. The Recruiter inserts the unit into a set of incomplete units. When a unit is completed, the Recruiter is notified and removes the unit from

the set and from the Accountant. The same happens if the unit is destroyed, which occurs if the trainer was destroyed.

Inserting and removing from the set is logarithmic to the size. The set is never expected to be very large.

7.3 Building Structures

Like training units, building structures is done in two steps: the structure is built and then it is constructed. Contrary to training, structures require both workers and a location to be built. Workers are the only units capable of building structures, and must move to the build location to do so. Terran workers must stay and construct structures until they are complete and Zerg workers are destroyed upon building, while Protoss require neither.

However, all Protoss structures must be in close proximity to a Pylon when built, and stops functioning without. The only structure exempt from is the depot and the pylon itself. This only becomes important in later stages of the game where Players might have to carefully manage their space. The agent is not expected to build enough structures for this to be an issue, so it is mostly disregarded. A Protoss structure is considered *powered* when in vicinity of a pylon.

Structure Placement

The Protoss player must place at least one pylon in any area he wishes to build in, and could distance pylons in a base to maximize coverage. Alternatively, if only a few pylons power a structure it creates a liability. The pylon could easily be destroyed to disable the structure, which is useful if it is a defensive structure or unit trainer. Fortunately, we expect the agent's bases to be compact enough such that we can place pylons somewhat arbitrarily and still manage to fit needed structures in the region. It is important however to place at least one pylon in a base before other structures can be built, but this is solved by the strategy modules.

Unless a building location is specified, the easiest solution is placing a structure as close as possible to the base location in the desired region. All the workers in the region are usually at the mineral fields by the depot, and therefore will not be far from the placement. The workers and depot become sheltered by

the structures, such that the opponents' troops are forced to move through bottlenecks to get the workers. Even if a region only has one exit, it might be desirable to spread structures as flying units can attack from any angle. Finally, the solution is easily implemented.

There are some exceptions. Depots will always be placed in new base locations and refineries can only be built on of vespene geysers. In both of these cases the building location will be specified by the superior module ordering the structure.

The agent attempts all locations in the map in a spiral pattern around the depot, and returns the first location that is available. This is simple, usually cheap but very expensive asymptotically, as it will be linear to the map size. However we never expect to visit anywhere near all the locations, since few tiles in the map cannot be built upon. To determine availability of a location, it must be clear of units and the tile terrain itself must be able to be built upon. Both of these are handled by BWAPI. Note that this solution places structures in a square pattern, which is not actually the closest to the base location except in *taxicap geometry*.

Placing structures too close can block passage between them, especially for larger units. The hitbox is arbitrarily different between otherwise equal sized structures in StarCraft, such that some combinations next to each other will block some units but not necessarily all. This is because the structure hitbox usually does not cover its location completely, allowing some leeway for smaller units.

By placing structures at least one tile from each other, we ensure all units can easily pass through. This is done by keeping a map of all owned structure locations, where their dimensions are increased by one tile in all directions. If a build location overlaps any of the occupied tiles it is determined as not available. Registering new structures and querying this is constant time operations, but the space used is linear to the map size. It could be improved by only keeping a map for regions we have bases in, although not an asymptotic improvement.

Additionally the agent avoids placing structures between resources and depots to avoid obstructing workers. By drawing the smallest rectangle including the depot location, geyser(s) and mineral fields, the agent avoids blocking workers by not placing structures within this. Building the rectangle is linear to the amount of items in it and querying it is constant, but we expect the amount of items to be low (ten or less). This structure map is handled by the *Base Manager* module, which the Architect is superior of.

To summarize, if a location is not specified, structures are placed in a spiral pattern around the depot, distanced by one tile and outside the gathering-zone.

Aquiring Builder

Recall the Task Master module from chapter 6. Since all workers are separated in base locations, we can easily pick a worker from the region the building location is. As the task master marks the jobs of all workers, we can pick a worker from either idle, mineral mining or gas harvesting (in that order). This way we do not interrupt other possibly important tasks. The agent searches through the groups until it finds a worker that is not carrying resources, which then becomes the builder and is tasked as such in the Task Master.

The time complexity is linear to the amount of workers in these groups. However, it is probable that less than half the workers are returning resources at any given moment, so it is expected that we only need to check two workers before finding a viable candidate. In case there are idle workers, the operation is constant.

This solution does not work if we have no workers in the region. This is the case when expanding, where the worker must be specified by a superior module.

Executing Command

Once the builder has been retrieved, it will be commanded to build at the specified location. A player cannot build in hidden terrain, so the worker will first be moved closer to the target location in this case. When the entire placement has been revealed the builder is commanded to build the structure.

The build order will then be stored in a set, containing the structure type, builder and location. Every frame, the Architect verifies the validity of the build orders, including verifying the builder has not been re-tasked. It also reissues commands to builders if required. As with monitoring trainees, incomplete structures are important to keep for the same reasons. In case of structures, there are more things that could go wrong which would incur cancellation of a scheduled structure. The builder might be destroyed or the build location might prove to be invalid upon being revealed.

All invalid orders are removed, forcing superior modules to reissue the build order. This is desired compared to repairing the order, as the build order might no longer be required.

Monitoring Constructions

Since Protoss structures auto-construct, the implementation is identical to monitoring trainees with the same time and space complexities.

As mentioned prior however, refineries are handled by the Morpher module. BWAPI is notified when a unit morphs, upon which it is inserted in a set of incomplete morphs. A morphing unit is not considered constructing, although it is incomplete. When a unit is finished morphing, no event is called, therefore forcing the agent to verify all morphs every frame. However, there are very few expected morphing units at any given moment.

7.4 Building Expansions

Recall that *expansions* are additional depots beyond the initial one, built at new resource clusters. As explained in the chapter 6, it is often profitable and sometimes necessary to expand resource harvesting to new locations. To avoid transporting cargo all the way between regions, players have to build depots near resources they wish to harvest.

When BWTA analyzes a map, it marks all viable depot locations. These are positions in which a depot will be at optimal distance from nearby minerals and geysers. Usually every resource cluster only has one of these positions, but some locations may contain overlapping depots. Regions rarely have more than one base location. Given a region, we can obtain all internal base locations, including these depot locations.

This significantly simplifies the construction of expansions, however the agent still has to decide which base location to settle.

The *Settler* module is responsible for all expansion behavior, and is also used to determine if expanding is possible. To build an expansion we need to specify a location and builder. These are handled by the Settler module, while the rest is done by the usual build procedure in the Architect.

Expansion Placement

To determine whether a base location is fit to expand it must not already be expanded upon or contain enemy forces. Additionally, it must be reachable

by non-flying units and have a path of regions without enemies. If these are satisfied, we consider the base location available for expansion.

From the starting region, we visit neighbor regions recursively in order of proximity by using a priority queue, where the lowest priority is picked first. Initially the start region with priority 0 is inserted. When we visit a region, we visit all contained base locations. If any are available, we return one as the target location of expansion. In case none are available, but the region is unoccupied by enemy forces, we add all unvisited neighbors to the queue. The priority is equal to the distance between the two region's centers plus the current regions priority. If the region does have enemy units within, its neighbors are not added to the queue. The algorithm terminates when the queue is empty or an available candidate has been found.

This solution satisfies our expansion requirements, and additionally includes the starting region should the depot be lost. It is very crudely implemented however, and quite expensive. Determining whether a region is occupied involves checking all enemy units. Given r regions and n enemy units, the time complexity is rn . This could be optimized by first marking all occupied regions before the recursion.

The enemy units are retrieved from the *Archivist* module, detailed in chapter 8.

There could exist more factors when determining expansion locations beyond distance. Resource quantity or defensibility are also important, but are not considered here. One could favor regions distant to the opponent for added defense or alternatively closer for better map control.

Expansion Builder

An expansion is often the first structure in a region, and therefore the player must acquire a builder elsewhere. The agent attempts to pick workers only from the region the location is within, so the Settler must specify a worker beforehand.

In this case, we just pick a free worker from the main base. There could be workers closer to the destination in other regions, and in this case the solution is sub-optimal. Recursively checking neighbors in order of proximity would result in an approximation of the closest worker. It was not a priority however as the improvement would be insignificant.

The solution is cheap, and costs the usual to pick a worker. A worker considered available is either idle or tasked with gathering.

Exploration

One of the challenges in StarCraft is the imperfect information players have of the world. It becomes necessary to deduce the opponent's strategy from a few units, and regular scouting is mandatory. The player with the best information coverage can make the most optimal decisions. This chapter describes how the agent solves the information and scouting problems.

All units, including buildings, have a line of sight equal to their *sight* value. Within this distance, the map is revealed to the players. Everywhere else, only the terrain features are visible or buildings in their last seen state. This is the *fog of war*, and all units within this are *invisible*. These units cannot be targeted or attacked, and the agent cannot access any data of them through BWAPI. It is up to internal data systems to store this information.

The entire map terrain, including initial resources, are revealed to players. All possible start locations are also known, but it is not known which the opponents occupy.

Ground units cannot see past cliffs or ramps, giving an edge to units on the high ground. Most start locations are located on plateaus, making them easily defended. Units or obstacles do not break line of sight however. Some units are *cloaked* or can *burrow* which renders them invisible to normal units even outside fog of war. BWAPI however does not limit the agent from accessing their data

while they are within line of sight. Units with the *detector* ability reveal cloaked units within its line of sight.

The agent must track data about discovered units, which is handled by the *Archivist* and *Geologist* module, the subject of the first section. The second describes the scout AI, implemented by the *Reconnoiter* module.

8.1 Tracking Units

Two things must be recorded for each enemy unit: position and type. The former is needed to track movements of the enemy army and to record enemy base locations. The latter is needed since the unit type contains all relevant values such as speed, damage and maximum health. Some units can change their type even, requiring regular updates of changes. The Siege Tank for example changes types when it goes into or out of siege mode.

The *Archivist* handles all known enemy units. The solution is sets of units and dictionaries of their data, where units are keys to their values. Every frame, the Archivist checks all stored units, updating their values if they are visible. Whenever a unit is first discovered or destroyed, BWAPI will notify the Archivist, which will then modify the dictionaries and sets appropriately. This solution spends logarithmic time on queries, insertions and deletions. The position and type queries could be improved to amortized constant time with hashing.

Units are inserted into multiple sets, separated in categories. Some units are stored in more than one set such that constant time retrieval is possible, at the cost of more marginally more space. These sets are used by other modules when specific kinds of units are desired, such as structures or workers.

The Archivist is one of the lowest in the hierarchy, as it depends on no other module and is used by almost all others. Other modules query the archivist for the units or their values.

Resource Locations

As all initial resources and their positions are retrievable from BWAPI at any time. Since new minerals are never created, it never becomes necessary to keep track of them here, as only the Gatherer module from chapter 6 needs them.

Vespene geysers however have special behavior in StarCraft, as mentioned in chapter 7. When a refinery is built upon a geyser, it is morphed rather than constructed. However, if the refinery is destroyed, a new, distinct geyser unit is created in its place. While all the initial geysers can be retrieved from BWAPI, new ones cannot. Otherwise, agents could deduce if any opponents lose refineries. This makes the tracking of geysers a bit more tedious than other units.

The *Geologist* keeps track of current known geysers in the map. They are stored in a dictionary of sets, where the keys are regions and the sets contain geysers from that region. This is because superior modules need the geysers separated in regions. The Geologist responds only to events, inserting new geysers or deleting refineries when discovery. All initial geysers are added on initialization. Since a refinery were the same unit as the late geyser, they can be used to remove them from the sets. Given r regions and g geysers, insertions and deletions cost $O(\log r + \log g)$ time. Since every region rarely has more than one geyser, the time is more like $O(\log r)$, which is the cost of a query. As usual, this could be improved with hashing to amortized constant time.

8.2 Scouting

Scouting is important throughout the entire duration of the game. Early on, players need to know where the opponent is and which opening they are using. Later on, observing the opponent's army size, unit types, base expansions and tech level is important to properly combat their strategies. The opponent's race is known unless they pick random, in which case this also needs to be scouted. There are no bots playing a random race however, so this is not relevant to us.

Recall that all possible player start locations are revealed to players. Scouting for the opponent base is then just visiting other start locations than ones own. Maps used in the competitive bot matches contain no more than four start locations. It could therefore be the case that a player has to scout up to two bases before knowing the enemy start location. Usually a player wants to know more than where the opponent starts, so even when it is deductible where the opponent is, the player wants to scout the base itself.

Once the enemy base has been revealed, and with it his opening strategy, the scout is not useful any longer. It is a long trip back to gathering resources and the scout could be followed, revealing your own location. Harassing enemy workers however can disrupt the opponent's economy, even if none of them are destroyed. Attacking enemy workers forces the opponent to pull at least one worker from gathering to defending. Usually there are no base defenses or

troops at this early in the game, except if the scout had to visit three bases. Scout harassment can be very advanced, as the scout can retreat while chased, possibly attacking passive workers. Killing even one worker is a big advantage this early in the game.

The *Reconnoiter* module contains the scouting AI. Using the TaskMaster module from chapter 6, the Reconnoiter picks some free worker as the scout. As usual, a worker is considered free if it is either idle or gathering resources. The scout is removed from the worker pool, as it will be a scout for the rest of its life span.

Usually a Protoss player will send the scout after the first pylon has been built, where the builder is used as a scout. As an approximation, the Reconnoiter will only acquire a scout if a specific supply total of 8 has been reached. This is because the first pylon is usually built around eight workers, such that the Reconnoiter should acquire a scout at the same time the pylon is built. The worker will never be picked however, as it is not considered free the moment the supply limit is reached.

If the Archivist has not recorded any enemy buildings and the Reconnoiter has no scout, it will attempt to acquire a scout. This implies scouting will proceed as early as possible and if all known enemy buildings has been destroyed, extending its use into later stages of the game.

While the Reconnoiter has a scout, it picks an unexplored base location and moves the scout towards there. A tile is considered unexplored by BWAPI if it has never been revealed for the duration of the game. Note that this removes the agents own starting location from the candidates. Once the scout reaches the destination the tile will be explored, removing it from the possible scouting locations. In case this is the opponent's starting location, their depot will be revealed as well.

There are alternative uses of the scout compared to harassing, such as building blocking the opponent's main base exit with defensive structures or building troop producers close by. These however must be planned in concordance to the grand strategy, as they involve spending resources. Some bots use these strategies, LetaBot for example can rush with bunkers.

At later stages in the game, faster units should be used to scout the map. In particular the Protoss Observer is a good scouting unit, as it is both permanently cloaked and a detector.

Combat

Combat in StarCraft can easily become the most complex part of the bot development. Numerous studies has been done on just small subsets of combat scenarios, where we are given specified units and enemies, but even these just scratch the surface. Like the macroscopic strategies in StarCraft, there is probably no optimal command scheme for ones units. Predicting the outcome of a combat scenario, is therefore only done as approximations.

This is also one of the areas where bots excel against humans, since they can easily command different units in complex ways.

The **ArmyManager** is the combat parallel to the **TaskMaster**. It contains all fighter units controlled by the agent and their current assigned **duty**, same as **tasks** for workers. A fighter unit is either idle, moving towards the opponent, attacking or defending.

9.1 Attacking

A player will almost certainly expand to their natural expansion first and a third base is only viable beyond early game. The natural expansion is very close to

the main base and is usually right outside the only exit of the main base. This proximity allows us to consider both bases as a single base. All this together we can assume the opponent has only one base for the duration of the early game, which is where the bot mainly operates.

In euclidean space the shortest path between two points is the same regardless of which is the origin. In terms of StarCraft, the shortest path to the enemy base, is also the shortest path from the opponent's base to yours. So assuming each player has only one base and both players will use optimal pathfinders, we can model the map as a straight line between the bases. Ours and their army will travel only along this path, and there is therefore no case of armies moving past each other without collision.

As a rushing bot, it is in our favor to move our troops as close to the opponents base without confrontation before they are ready. This pressures the opponent while also providing some scouting, and allows us to attack as quickly as possible. When we predict a victory in combat against the opponent army, we attack. In some cases it might be prudent to wait until more units have been amassed, especially if we produce more troops than the opponent, as we will sustain fewer losses and be more certain of a victory. This is difficult to asses while also being very risky, so the safer option is just attacking immediately.

Attacking is handled solely by the **Attacker** manager. Predictions are handled by the **CombatJudge** module, which given a set of units outputs their strength value.

Prediction

Predicting combat is very difficult in StarCraft. Even beyond the numerous factors in combat, the optimal command of ones troops is very Dependant on the command of the opponents'. Unless the opponent operates in a recognizable pattern, they movement is unpredictable. Predicting combat is therefore more based on what units the opponent controls, how the terrain is and then either a theoretical upper limit of their damage output versus a prediction of our own.

Some of the most successful predictors simulates a simple form of combat and evaluates the result. This however assumes how the opponent will control its troops, or at least assumes it knows the optimal control scheme. Neither of these are possible, especially not with incomplete map information, but it has proven to be close enough to the correct result.

It is assumed the opponent has only a single army which is always gathered in

close proximity. This is not very useful in late game where flanking maneuvers and harassment is viable tactics, but it is sufficient for early game. The worst case is we will overestimate the opponent strength, which is a much better case than underestimating it. It would be possible to detect these individual groupings however with a cluster detection algorithm.

TODO Prediction heuristic.

Targeting

Usually RTS games has a command called *attack-move*, and StarCraft is no exception. A unit executing this command will move towards a target or location, but will attack any enemies along the way. In StarCraft the unit will prioritize units that are attacking it before others. So beyond auto-targeting nearby enemies, this command also carries a simple prioritized targeting.

By targeting the enemy base location or structures with this command, we have simple prioritized targeting. We avoid targeting the opponents units, as our army might be lead astray by a decoy. The most important targets are the enemy structures since a player loses when they have no more of them.

Since we already have a dictionary of enemy buildings, the agent just picks one arbitrarily from the list. Assuming the opponent has only one base, the army will attack the foremost building regardless of target since they are attack-moving. Since the natural extension is usually the only exit from the main base, this assumption is valid while the opponent has two or fewer bases, which is true for the entire early game.

Some bots go beyond the attack-move prioritize. A useful alternative is attacking the unit with highest health to resources cost ratio, targeting damaging units first. This will ensure an attack deals a size-able blow to the opponent's economy, trading resource loss as favorably as greedily possible. A similar one is attacking units with highest health to damage ratio. There is a lot more to prioritization. This is especially true when considering large scale armies composed of different units, where counter-play between individual types are important. Beyond fighting the enemy army, the order when attacking is usually opponent's army, then economy and then unit facilities.

Troop Rendezvous

As all troops are sent off immediately from production, its necessary to avoid sending them into the enemy base one at a time. The strength of an army increases faster than linearly with its size, as troops will die slower while enemies die faster. This means our troops get more time to deal damage while their troops get less.

At some point troops in transit will be within some danger distance of the enemy base, where they will risk being attacked. This will be the minimum distance from the opponent where they can safely gather. By only counting these arrived troops in our combat predictions, we can tell when enough troops have gathered to attack. Currently, the bot only counts units currently within this danger distance of an enemy unit as arrived, that is, it does not record arrivals. Testing proved the battlefield shifted too much to record arrivals and pausing them until attack, which lead to troops being scattered across the battlefield rather than actually gathering. This solution proved sufficient however for the few and small units the bot uses. A better solution would be needed with a larger army composed of larger units. This could be solved with a cluster detection system, to detect when ones army is actually gathered.

Fighting

There is a lot of depth in unit combat in StarCraft. Maneuvers and strategies involving the precise commanding of units is called *micro*. There has been multiple studies in this area in specific scenarios.

The bot uses the Zealot unit which is a very efficient unit but also a very simple unit. All we need to do attack move each zealot, which will then auto-target to the nearest enemy. This is surprisingly effective against other units, mostly because other units are a lot more complicated, especially if they are ranged. There is a lot of room for improvement, but most of the bot development focused on the macroscopic strategy than the micro.

Retreating

Sometimes because of incomplete information, the combat prediction turns out to be wrong and the opponent had in fact more units hiding the fog of war. In some cases, it is optimal to retreat. Depending on how tangled ones units are

with the opponents, retreating becomes a less desirable option as some units will die without fighting back. However the opponent would not lose very much value in troops by staying in combat, compared to how much you would lose by retreating, then it is the better option to retreat. Calculating this is not easy however, and would also require a strong combat predictor, something that we don't have access to.

The bot will therefore not retreat any units that has been assigned combat, only units that are in transit. The logic here is that we don't know how tangled our combat units are to the opponent, but the transit units will usually be outside enemy range, such that we lose no troops to retreat.

Retreating units move towards the start location, which will path them outside the enemy base. This is usually the optimal solution barring cloaked or flying units, or enemy units that need to be avoided. Fortunately, its rare that the opponent has some troops that can be avoided if they are on the retreat path, and the bot does not utilize cloaked or flying units. A solution to these problems however would be using a custom pathfinder using heatmaps of enemies or potential fields.

9.2 Defending

Sometimes the enemy attack first, break through our attack or somehow avoid our army. In this case the base must be defended. At other time, the opponent might have sent an early Zerg rush or a worker as scout and harasser. In these cases we need to pull workers to fend off the attackers.

Although they have not been implemented, turrets are effective when playing defensively. Usually cost-effective compared to mobile units, they can easily repel early assaults. The problem however is finding proper locations, as the bot has to predict the point of entry to defend it. This is coupled with the logistical problem of blocking the entrance or bottlenecking it.

The **Defender** manager handles the enlistment and commanding of defenders.

Scrambling Defenders

The problem with designating defenders is related to retreating. We must avoid drawing troops away, that would be better used attacking. This is difficult to

asses. It is obvious however, that everyone in a region under attack, should defend it. This will not interrupt desired combat or cause undesirable behavior. If they are already defending, nothing is interrupted, and otherwise they are better served defending.

As it turns out from testing and because of the straight-line assumption, we very rarely have any units left beyond our base if the opponent is attacking, since they would already have fought and died. Those that are beyond however, continue to assault the opponents' now undefended base, which is desirable since they are too far away to actually help the defense. The current simple implementation therefore proved to be sufficient.

Militia

Sometimes there are not enough troops in the base. This usually happens before an army has been acquired or after it has died. In this case, a common strategy is scrambling all the workers for defense in hopes of surviving. Worst case all workers die and the player loses, but that would have been certain without the militia anyways.

Workers are tasked to defend until `CombatJudge` predicts they win. If its an early attack, usually a few workers, but later in the game it is almost all workers.

CHAPTER 10

Macro Strategy

This chapter concerns itself with how we implement the overall strategic reasoning in our bot. We have already gone through the basics of strategies in StarCraft, and all the basic functionalities has been implemented in the bot. Now it must all be put together in a cohesive form to achieve a strong AI.

Initially the bot functioned purely on a greedy principle. Concerning a specific set of units, we attempt to produce them if possible every frame. This is possibly the fastest and simplest AI to implement, but it has proven to be an effective approximation in some algorithmic fields. As the bot was developed additional systems was implemented, but the core auto-pilot remained greedy.

The overall strategy of the bot is rushing. It needs to be effective in economy early on and build troops as fast as possible. It sacrifices long term advantages for an early game attack. If the initial attacks are successful, the bot would gain even stronger long term advantages by hindering the opponent, and would be able to transcend into mid-game tactics.

While one could focus on a niche strategy such as proxies or such, as explained earlier these strategies are strongly countered by some specific strategies. They are high-risk, high-reward, but only work in a small subset of match-ups. So even though they would be effective in these match-ups, we would need to implement other solutions for the rest, where a not insignificant amount of them

would require an advanced bot beyond these 'cheese' strategies. Therefore, we could just as well implement an advanced bot and cover all match-ups. All-in would be such a niche-strategy against humans, but bots have proved to be very ineffective against this strategy. Additionally, the all-in strategy is a more aggressive version of rushing, so it could easily be transformed into the latter without any large scale implementations.

All combat decisions are made by the **Strategist** manager, while economy decisions are made by the **Economist** manager. These are disabled by the **Despot** manager when executing a build-order. The **Despot** module is the highest in the hierarchy, beyond the core AI module receiving call-backs from StarCraft.

10.1 Economy

TODO intro

The **Economist** manager moderates worker production and expansions, unless it has been disabled by the **Despot**.

Workers

Generally it is advised that a Protoss player keeps producing workers constantly until late-game, but some openings might temporarily stop worker production. Therefore, if we are not at the desired worker amount in a base, we attempt to produce a new worker with regards to resources and whether a worker is already in production.

Expanding

An expansion secures a more permanent economic advantage, as a player secures both more resources and gathers them faster. More expansions also secures map-control, as the player will have more presence in the map and control more resources.

Deciding when to expand is both difficult and pivotal. It takes time and resources and leaves the new base pretty vulnerable unless there is a standing

army. A player with a stronger army can safely expand while guarding it, securing his advantage or at least keeping up with the opponents expansions. Without the stronger army its a risky maneuver, but might still be viable or even imperative. At some point the resources in ones current bases will be depleted, leading to a need to expand. If a player expands before having saturated his current mineral fields, its called a *fast expand*, sometimes used as an opening. This is very risky but with a very immediate payoff in economy. The contrary formula, keeping a single base until mid-game or alike is called *one base play*.

Implementation

It is only clear that an expansion is viable when all mineral fields have been saturated. Before this, it is both risky as we sacrifice current resources for later economy, which requires both a strategic reason and foresight whether it is viable. By expanding when we have saturated our mineral fields, we risk expanding too late, but ensure a consistent worker production throughout the game. Expanding too late is better than expanding too early; an early expansion will waste resources or possibly result in an immediate loss, where a late expansion simply wastes an opportunity.

Since expanding is a costly affair, the bot needs to save up resources. Therefore, the expansion is queued in the build-order, thereby pausing other productions until the expansion is ready.

TBD Implement better expanding, avoid opponent regions and avoid if in combat.

The bot can perform a fast expand opening, but it has not been very successful compared to its usual one base play. The difficulty lies in detecting whether a fast expand is viable, which is very much Dependant on enemy openings. It was decided to focus on the safer opening instead as it is a more versatile strategy.

TBD Evaluation of the auto-expander vs. one-base bot.

10.2 Combat

TODO Combat intro.

TODO Building troops.

TODO Counter Strategies.

Strategist order the production of troops and troop facilities.

Troops

As this bot is very simple in terms of handling units, it is appropriate to use the first available units and attack as quickly as possible. The idea behind this strategy is, that since this bot lacks greater strategic reasoning and handles late game units poorly, then it should attempt to win the game before any of that becomes relevant. This can beat more complex bots, simply because their otherwise superior strategies are never allowed to hatch. Additionally, Protoss has the strongest and easiest to use early game units, and when massed will easily beat both Terran and Zerg.

Currently the bot just mass produces zealots whenever it can in greedy fashion. As producing zealots is the ultimate goal of the bot, this is sufficient.

TBD Additions: Dragoons, upgrades.

Facilities

It was found through testing that a single saturated start base can handle around four gateways constantly producing zealots. This gets more difficult to asses with expansions, and it was unfortunately also found through testing that greedily building limitless gateways is not even close to optimal.

The current limited solution is building seven gateways at max. Usually if the bot reaches a point where it needs more, it has probably already lost due to lacking in technology. A scalable solution has not been implemented, but to be optimal it would require some estimation of resource input, along with estimated need for profit for further development of base and technology.

10.3 Build-orders

Openings in StarCraft are usually detailed like a *build-order*, which is very much like a chess opening. Because of the huge decision space, even humans need

some guidelines weathered by experience to perform well in the beginning of the game. The optimal build-orders have changed a lot across the years, shaping the meta-game of StarCraft. As new strategies were found, new openings had to accommodate new possible opponent strategies.

A build-order is a list of units that must be produced in a specific order, sometimes after certain events. There is no timer in vanilla StarCraft and using one is considered cheating in tournaments, so items in the build-order are to be constructed when specific supply costs has been reached. Between each item in the build-order, the player must produce workers until then next supply limit has been reached unless specifically otherwise instructed.

This is akin to executing a queue of orders. One could make a queue where the next item is executed if a specific supply is reached while automatically building workers like a human player. The easier solution for bot however is of course manually filling in the gaps with worker orders, such that the supply limits will be reached at the correct times (assuming no workers are lost). Therefore, the solution in the bot is simply having a queue of units. While the queue is not empty, we do not execute any higher-order manager AI that would otherwise automate production. When the queue becomes empty, we resume automatic production.

The benefit of this solution is the ease of adding new openings to the bot. Although the general autopilot AI is the same, different openings will vastly change the strategy and outcome of the game. Proper openings are key to winning in early game, or at least not lose during. The downside however is the rigid planning structure, which can't allow for easy adaptations or responses. What if the bot loses a worker or a structure? It is possible that the current build-order is no longer optimal, viable even. We avoid this, by making sure the build orders are short, such that the opponent could in no way interrupt with any units other than harassing workers.

Additionally, this solution could be used to pause other production in mid-game. Since the bot operates in a greedy fashion, it can be difficult to build high-cost units. The solution is to enqueue the unit, pausing all other production until the unit has been created. This is used in the agent when expanding, where the depot is queued. This solution is risky however, as the bot will not consider removing the depot from the queue under any circumstance. It does not seem to cause any problems however, as the bot will reach the amount of resources required very quickly, resuming normal production again.

Since the build-order is a queue, executing it is constant time, and enqueueing the build-order is linear. The build-order is never changed during gameplay.

The **Planner** module handles executing build-orders. The main module **Despot** handles queuing build-orders and disables other managers while **Planner** is executing a build-order.

Results

Across 780 games the one-base version garnered 527 wins, netting a 79.23% win-rate in a mixed-division on the SSCAIT tournament page. Unfortunately there is not a separate statistic for student-division only, but it was usually contesting the UAlberta bot for top student bot. The UAlberta bot usually wins against this bot however.

TBD always expand bot + condition expand bot results.

As one could expect, the bot does well against others using boom-strategies and proxy-strategies. It wins by far most of the matches against other rush bots, including mirror-matches, except UAlberta bot as mentioned. It fares poorly against the turtling Terran bots, which quickly produce more advance units and win through technology. The final additions to the bot were attempts to combat the turtle strategy, however it was never effective enough. It is the natural counter to the rush strategy.

Performance wise, the bot has not broken any of the SSCAIT time limits. Recall that they were:

- No more than 1 frame longer than 10 seconds.
- No more than 10 frames longer than 1 second.

- No more than 320 frames longer than 55 milliseconds.

In fact, it has not been observed that even one frame exceeded 55 milliseconds. This is not unexpected, since the bot has favored cheap and simple algorithms.

The bot did not reach a very complex higher-level AI. The focus was creating a competitive bot, so development focused on areas that improved performance the best. This is expectedly in low-level technical features such as using the command interface and handling units. The opening-moves and attack commanding is the most advanced AI the bot can boast. While as a whole it performs somewhat intelligent, the lack of adaption means it is really not.

Even though the bot has now grasped the basics of StarCraft and RTS gameplay, there is probably a long way yet before any high-level AI implementations would be effective. Its still does not use most of the Protoss units and buildings and does not research upgrades. The next steps to take would be improving its combat predictions and map awareness with regards to enemy unit placement, such that it could understand separate bases and armies. From there, counter strategies would be necessary, especially against flying and cloaked.

As an additional note, the bot has problems pathfinding in the maps *Circuit Breakers (4)* and *Heartbreak Ridge (2)*, where neutral structures block terrain passages. The bot uses the built-in pathfinding, which cannot path around these kinds of obstacles. In *Circuit Breakers*, this will result in a loss every fourth game, as it is only an issue in one of the four start locations. In *Heartbreak Ridge*, it prohibits the bot from expanding to a certain base location. This problem was never fixed as it would require the implementation of a custom pathfinder which would be too time consuming. Interestingly, the built-in bots suffer the same issue.

The bot will be applied to at least to the 2015 AIIDE tournament with a submission deadline the 15th of august. There is however a limit to 30 submissions, so it is not certain the bot will appear in the tournament.

CHAPTER 12

Conclusion

The bot has been a success with a 79.23% win-rate as Protoss in mixed-division on SSCAIT. It handles basic StarCraft maneuvers such as harvesting both types of resources on multiple bases, building structures and executing a simple rush strategy. It can execute build-orders, including early expansions. When in combat, it can do basic tactical decisions such as gathering troops and retreating. With a very simple but somewhat effective combat prediction heuristic, it will attack only when ready. The project has achieved its goals of developing a competitive bot for StarCraft. While the bot is simple compared to humans and older bots, it is still above in complexity and effectiveness than the average.

The bot will be submitted to the 2015 AIIDE competition, and will probably be further developed at least in this year in preparation for later competitions.

Bibliography
